

README

Athell Felix Ikechukwu O. Dec 31, 2023,
Dynamic arrays: `std::vector`

Introduction to containers and arrays

we be talking of vectors now.

The variable scalability challenge

Due to the fact that we would eventually interact with large quantity of data in some scenerio, it can prove cumbersome handling them (due to thier leviathan sized nature) without falling victim to error.

c++ has a solution to the problem: `containers`

Containers

think of containers as egg crates: they provide an easier, efficient and portable way of handling as many eggs as the crate is designed to hold.

containers in c++ exist with the same purpose as egg crates (in a way). they make it easier to manage (potentially) large collection of objects (called `elements`).

a `string` is a subtle example of containers in c++: they provide a collection of characters, which can be outputted as text.

The elements of a container are unnamed

while the container object itself has a name (else how would we refer to it?) its elements are unnamed: this is an `important` details of containers.

containers provide interfaces to enable access to the its elements.

The length of a container

teh number of elements in a container, in programming generally, is known as its `length` , or sometimes count.

```
std::string name { "Athell" };  
std::cout << name << " has " << name.length() << " characters.\n";    //  
here the length member function is used to return the length of the  
container: std::string name.
```

Container operations

containers should basically perform tasks:

1. create an container (empty or initialized)
2. access to elements
3. insert or removal of elements
4. get the number of elements in the container.

any other operations, that should be beneficial to the logic of the container.

Element types

in most programming languages containers are `homogenous`, meaning the elements of the container are required to have the same type.

some containers use a preser type for its elements: like with `std::string` its elements are of type: `char`.

in c++ containers are usually made as class templates, to make it possible to have containers for any specified datatype and length: this makes it really convenient.

Introduction to arrays

an array is a container data-type that stored its elements `contiguously` (meaning everly element is placed in adjacent memory location with no gaps).

allowing fast and direct access to any element.

c++ generally contain 3 primary array types:

1. C-styled array : known to behave strangely and are dangerous, wth did you expect? its from fucking C.
2. `std::vector` container class
3. `std::array` container class

Introduction to `std::vector` and list constructors

here we be talking of `std::vectors`.

Introduction to `std::vector`

the `std::vector` is one of the containers in c++ library that implements an array. it is declared in the vector header as a class template, with a template parameter

that defines the type of its elements. therefore a `std::vector<int>` defines a container with elements of type `int`.

```
#include <vector>

int main()
{
    std::vector<int> empty {};
```

Initializing a `std::vector` with a list of values

most often we want to initialize our containerw with values: we can achieve this using list initialization and C++17 CTAD (Class Template Argument Deduction)

```
#include <vector>

int main()
{
    std::vector<int> arr_int { 1, 2, 3, 4 ,5 };           // explicit delcaraion
    of a container of int elements
    std::vetcor arr_char { 'a', 'b', 'c', 'f' };         // initializing and
    determining its type using CTAD.
```

List constructors and initializer lists

containers have a special constructor known as `list constructor`. it allows us to construct an instance of the container using an initializers list.

it does three things

1. ensures container has enough space to hold all the initialization values
2. sets the length of the container to the number of elements inputed
3. initializes the elemenst in sequential order as listed in the list initializer.

Accessing array elements using the subscript operator (`operator[]`)

just like in C: they are `zero based` tho.

an example:

```

#include <iostream>
#include <vector>

int main()
{
    std::vector numbers { 1, 2, 3, 4, 5 };
    std::cout << "the first number is " << numbers[0] << '\n';
    // `[x]` is the subscript operator with x being the subscript: the
    // index to be accessed.
    std::cout << "the second number is " << numbers[1] << '\n';
}

```

Subscript out of bounds

when indexing an array: the index selected must be a valid element in the array, its subscript limit is usually $N-1$ where N is the size length of the array.

the `operator[]` does not perform any sort of bound checking: meaning passing in an invalid index will lead to UB.

Arrays are contiguous in memory

one defining characteristics of arrays is that they are allocated contiguously in memory: meaning they are located adjacent in memory with no gaps.

an example:

```

#include <iostream>
#include <vector>

int main()
{
    std::vector primes { 2, 3, 5, 7, 11 };
    std::cout << "an int of of size " << sizeof(int) << '\n';
    std::cout << &(primes[0]) << '\n';
    std::cout << &(primes[1]) << '\n';
    std::cout << &(primes[2]) << '\n';
}

```

the results on my system:

```

an int of of size 4
0x55e84ecf3eb0                                     # if you
look closely the address increment with 4 proving they are contiguous.

```

```
0x55e84ecf3eb4
0x55e84ecf3eb8
```

arrays are one of the containers that allow `random access` meaning every element can be accessed directly with equal speed, regardless the length of the array.

Constructing a `std::vector` of a specific length

`std::vector` has an explicit constructor `explicit std::vector<T>(int)` that takes in only one integer value to set the length `std::vector` when constructing.

an example:

```
std::vector<int> data (10); // we just declared
a vector of an immutable length: 10.
```

each of the elements are value-initialized and are therefore all initialized to zero.

this constructor must be called using `direct` initialization not `list`.

List constructors take precedence over other constructors

yup read this in the site. moral lesson is: list constructors take precedence to other constructors and that's why initializing with direct initialization will only set the length of an array and zero initialize its members and using list will not (it just provides the elements).

note:

```
std::vector<int> num({ 10 }); // this is an alternative version of
copy list initialization so this will not construct for specified length.
```

Declaring vectors as a member of a class type

due to the fact the direct initialization is not allowed for default member initializers: we can't do this to set the length of a vector:

```
struct Foo
{
    std::vector<int> vec ( 10 ); // error due to stated above
    reason
};
```

but we can do this instead: creates an anonymous object and then use that to defaultly set the length of the vector via the copy constructor

```
struct foo
{
    std::vector<int> vec { std::vector<int>(10) }
}
```

Const and constexpr std::vector

they can be made const and this makes their elements immutable: also they must be initialized:

```
const std::vector<int> vec {1, 2, 3, 4};
vec[0] = 9; // error they are
immutable.
```

also normal std::vector elements are not allowed to be const: they must be mutable.
so doing something like this is prohibited.

```
std::vector<const int> vec { 1, 2, 3, 4 }; // illegal: error:
`<const int>` can't make the int const cause that will make the elements
immutable which is not allowed
// vectors.
```

std::vector and the unsigned length and subscript problem

read this section and reread the paper bjoerne stroustrup wrote regarding the mistake of using unsigned integers for subscripting due to it bound wrapping into garbage values because it relies on modular arithmetics

Paper: [why subscripts should be signed](#)

A review: sign conversions are narrowing conversions, except when constexpr

unsigned int to signed int or the other way around are narrow conversions, which list initialization do not allow

but when made constexpr they allow them because the compiler would have to make sure the conversion is safer (at compiler time) otherwise halt the compilation

so narrow converting like declarations are allowed if the variable to be narrowed is a

`constexpr`.

```
#include <iostream>

void foo(unsigned int)
{}

int main()
{
    constexpr int s { 5 }; // now constexpr
    [[maybe_unused]] unsigned int u { s }; // ok: x is constexpr and can be
    converted safely, not a narrowing conversion
    foo(s); // ok: x is constexpr and can be
    converted safely, not a narrowing conversion

    return 0;
}
```

copy initialization do allow narrow conversions tho.

also: the non narrowing conversion of `constexpr int` to `constexpr std::size_t` is something we will use a lot. according to what i just read.

The length and indices of `std::vector` have type `size_type`

yup. and `std::size_t` is an alias for an `unsigned long` or `unsigned long long`.

each of the STL container classes define a nested typedef of `size_type`, which is the alias for they type used for the containers length or (indices if supported).

almost always the `size_type` is usually an alias for `std::size_t`, but it can be overridden (rarely done).

when accessing the `size_t` member of a container class er must scope the class with its full template declararion

an example:

```
std::vector<int>::size_type
```

Getting the length of a `std::vector` using the `size()` member function or `std::size()`

an example:

```

#include <iostream>
#include <vector>

int main()
{
    std::vector<int> vec { 1, 2 , 3, 4, 5 };
    std::cout << "the size of vec is " << vec.size() << " used .size()\n";
    std::cout << "the size of vec is " << std::size(vec) << " used
std::size()\n";
}

```

in c++ most containers have their length member functions dubbed as size, except in strings and string_view that have both length and size.

note:

- favor using `std::size()` that containers `.size` this is because it works on a broader set of containers like non-decayed c-styled arrays.

if we wasn't to use the `std::size()` or `.size` to initialize some variable with the size of a container, we might sometimes run into signed/unsigned conversion warnings or errors. we can prevent this by `static_cast`ing to the type same with the variable to be stored into. an example:

```

#include <vector>

int main()
{
    std::vector<int> vec {1, 2, 3, 4, 5};
    int length { static_cast<int>( vec.size() ) }; // error
    free!
}

```

Getting the length of a std::vector using std::ssize() C++20

it is a function that returns the length of a container as a large signed integral type (usually `std::ptrdiff_t`), it is used as a signed counterpart of `std::size_t`. if we want to use this as a method for assigning the length of a container to an `int` we have some options:

- first: using the `static_cast` to prevent narrow conversions:


```
std::vector num { 1, 2, 3, 4 };
int length { static_cast<int>( std::ssize(num) ) };
```

- or we can use the auto keyword to deduce the type of the length to be returned by the integral type:

```
std::vector num { 1, 2, 3, 4 };
auto length { std::ssize(num) };
```

Accessing array elements using operator[] does no bounds checking

yeah it doesnt if indexing for an out of bound index, UB can occur.

Accessing array elements using the at() member function does runtime bounds checking

yeah it performs bounds checking at runtime so if an out of bound index is requested using the `at()` member function it halts the program at runtime.

an example:

```
#include <iostream>
#include <vector>

int main()
{
    std::vector vec { 1, 2, 3, 4, 5 };
    std::cout << vec.at(4);
    // this is fine. still in bound.
    std::cout << vec.at(8);
    // terminates program at runtime: not in bounds.
}
```

when the `.at()` member function tries to access an out of bound index it throws an `std::out_of_range` error and terminates the program.

because of its runtime evaluation nature it's slower than `operator[]` indexing but more safer.

Indexing std::vector with a constexpr signed int

we can do this to enable the compiler convert the constexpr signed int to a `std::size_t` without it being a narrow conversion.

an example:

```
std::vector<int> num { 1, 2, 3, 4 };
std::cout << num[2];

constexpr int index { 3 };
std::cout << num[ index ];      // not a narrow conversion.
```

dont get the use of this but okay.

Indexing std::vector with a non-constexpr value

the subscripts used to index an array can be non const.

an example:

```
std::vector num { 1, 2, 3, 4 };
std::size_t index { 2 };

std::cout << num[index];      // operator[] requires an index
of type std::size_t so no narrow conversion required.
```

whn our subscripts are non-constexpr signed integer we run into some issues.

an example:

```
std::vector num { 1, 2, 4, 5 };
int index { 3 };
std::cout << num[index];      // hers a problem occurs:
narrow conversion (implicilty converted signed int to std::size_t).
```

this might raise some warnings (not errors) if warning flags where set, this can be quite annoying.

we can prevent this by always:

1. `static_cast` ing our subscripts when not of the right type which is `std::size_t`
2. or just always using the `std::size_t` as the type of subscripts whenever.

Passing and returning std::vector, and an introduction to move semantics

the `std::vector` element type is part of the type information of the object. so when we declare `std::vectors` as parameter ina function.

we declare them with their full template declaration. and because vectors can be quite expensive to copy we pass them as references.

an example:

```
#include <iostream>
#include <vector>

void passByRef(std::vector<int>& arr)           // we must explicitly
specify the <int>.
{
    std::cout << arr[0] << '\n';
}

int main()
{
    std::vector vec { 1, 2, 3, 4, 5 };
    passByRef(vec);
    return 0;
}
```

Passing std::vector of different element types

because the passByRef function only takes in vectors of element type `int` we cannot pass in another vector of a different element type.

in c++17+ this can be bypassed by using CTAD (class template argument deduction) an example:

```
#include <iostream>
#include <vector>

void passByRef(const std::vector& arr)         // no need to specify
the <int> to allow CTAD work.
{
    std::cout << arr[0] << '\n';
}

int main()
{
    std::vector prime { 1, 2, 3, 4, 5 };
    passByRef(prime);                          // CTAD uses its argument
to infer that the vector should be of type std::vector<int>
}
```

we can also make the function a template function instead of relying on CTAD:
an example:

```
#include <iostream>
#include <vector>

template <typename T>
void passByRef(const std::vector<T>& arr)
{
    std::cout << arr[0];
}
```

Passing a std::vector using a generic template or abbreviated function template

pass

Asserting on array length

when we try access an index that's out of bounds of a vector it leads to UB.
so we can assert on the size first: if logically acceptable we can proceed.

Returning a std::vector

when we need to pass in a vector into function we pass it as a const reference just not to make a copy.

but for some god-damned reason it's okay to return it by value.

Introduction to move semantics

when initializing or assigning a value to an object copies value from another object we say it's operating with `copy semantics`

sometimes when we initialize `some` objects with another objects of the same datatype instead of accomplishing our goal with `copy semantics` the compiler performs `move semantics` instead.

this is when instead of copying the data from an object to initialize another object, we move the data of the initializing object to the one to be initialized (`move semantics`)

there are 2 criteria for this to occur:

1. the datatype must support move semantics

2. the object to be initialized must be initialized with another object of the same datatype that's an rvalue.

not many datatypes support, `std::vector` and `std::string` support move semantics tho.

an example of move semantics in c++:

```
#include <vector>

std::vector generateVec()
{
    std::vector vec { 1, 2, 3, 4, 5 };
    return vec;
}

int main()
{
    std::vector vec { generateVec() };           // here move semantics occur
    // because firstly std::vector supports it and
    // secondly its being
    // initialized with an rvalue object returned by the generateVec() function.
}
```

We can return move-capable types like `std::vector` by value

Because of `std::vector` supports move semantics we can return by value without worrying about object duplication.

making return by value extremely inexpensive for this types.

Arrays and loops

an example:

```
#include <iostream>
#include <vector>

int main()
{
    std::vector scores {23, 455, 65, 76, 23};
    std::size_t length { std::size(scores) };

    int average { 0 };
    for (std::size_t index { 0 }; index < length; ++index)
```

```

        average += scores[index];
    average /= static_cast<int>( std::size(scores) );

    std::cout << "average score is " << average << '\n';
    return 0;
}

```

accessing the elements of a container in some order is known as `traversal` or `iterating` through.

Templates, arrays, and loops unlock scalability

an example:

```

#include <iostream>
#include <vector>

template <typename T>
T average(const std::vector<T>& array)
{
    std::size_t len { array.size() };

    T average { 0 };
    for (std::size_t index { 0 }; index < len; ++index)
        average += array[index];
    average /= static_cast<int>(len);

    return average;
}

int main()
{
    std::vector class1 { 84, 92, 76, 81, 56 };
    std::cout << "The class 1 average is: " << calculateAverage(class1) <<
'\n';

    std::vector class2 { 93.2, 88.6, 64.2, 81.0 };
    std::cout << "The class 2 average is: " << calculateAverage(class2) <<
'\n';

    return 0;
}

```

Arrays and off-by-one errors

fuck this shi

Arrays, loops, and sign challenge solutions

we all know that container subscript, index and length type are `std::size_t` which is an unsigned integral.

because of the unstable nature of unsigned types when interacting with negative values: this can pose some issues like so:

```
#include <iostream>
#include <vector>

int main()
{
    std::vector vec { 1, 2, 3, 4, 5 };

    for (std::size_t index { vec.size() -1 }; index >= 0; --index)
        std::cout << vec[index] << ' ';
}
```

firstly it iterates and prints the vector elements in reversal as expected then after (and at) the iteration where index becomes -1 undefined behaviour occurs.

this is due to the bound wrapping action unsigneds take when interacting with negative values.

because it bounds to some very large garbage value whenever

`index` should be negative the loop never ends (due to its `index >= 0` condition) and prints garbage values infinitely until the program crashes.

although there are some ways to fix this issue. using signed integral type can easily fix this but it itself has its own downsides.

an example:

```
#include <iostream>
#include <vector>

int main()
{
    std::vector vec { 1, 2, 3, 4, 5, 6, 7, 8, 9 };

    for (int index {static_cast<int>( vec.size()-1 )}; index >= 0; --len)
        std::cout << vec[static_cast<std::size_t>(index)] << ' ';
}
```

in this snippet although it works the cluttering of code is crazy: which reduces readability. although the cluttering is not that severe in this example it becomes really much in larger code. so yeah `unsigned` compromises execution safety when interacting with negative but `signed` creates a lot of clutter. so what can we do?

Leave signed/unsigned conversion warnings off

makes debugging easier and prevents being drowned with a buttload of warnings.

Using an unsigned loop variable

some developers prefer using this paradigm but we have to be careful to prevent running into signed/unsigned mismatches when doing so.

1. using the `size_type` of containers (usually it's just an alias of `std::size_t` which is just an unsigned long (long))
an example:

```
std::vector arr { 1, 2, 3, 4, 5 };

for (std::vector<int>::size_type index { 0 }; index < arr.size(); ++arr)
    std::cout << arr[index];
```

however using `size_type` has some downsides: it's nested so it clutters code and when used in a function template, because it becomes a nested dependent type (due to being used in a template function that takes in template parameter to generate the type of vector container thus making it dependent) we have to prefix the `typename` keyword when declaring it. you have to use the prefix `typename` whenever declaring a dependent type
an example:

```
#include <iostream>
#include <vector>

template <typename T>
void printArray(std::vector<T>& arr)
{
    for (typename std::vector<T>::size_type index { 0 }; index <
arr.size(); ++index)
        std::cout << arr[index];
}
```



```
int main()
{
    printArray( std::vector {1, 3, 4, 5, 6, 7} );
}
```

any name that depends on a type containing a template parameter is called a `dependent name`.

and they must be prefixed with the `typename` keyword.

sometimes some developers use alias to make loops easier to read.

an example:

```
using arrayi = std::vector<int>;
for (arrayi::size_type index { 0 }; index < arr.size(); ++index)
```

2. We can also tell the compiler to return the type of the container using the `decltype` keyword, which returns the type of its parameter.

an example:

```
for (decltype(arr)::size_type index { 0 }; index < arr.size(); ++index)
```

however if `arr` is a reference the above doesn't work so we have to remove the reference: and to do that we:

```
template <typename T>
void printArray(std::vector<int>& arr)
{
    for ( typename std::remove_reference_t<decltype(arr)>::size_type
index {0}; index < arr.size(); ++index )
        std::cout << arr[index];
}
```

unfortunately this is neither concise or easy to remember so most times programmers just use `std::size_t` because most of the times containers `size_type` are usually aliases of `std::size_t`.

Using a signed loop variable

What signed type should we use?

1. with small arrays `int`

2. with very larger arrays, or if wanting to be a bit defensive we use the strangely named `std::ptrdiff_t`. its a typedef of the positive counterpart to `std::size_t`.
3. because it has a wierd name long name, another approach is using an alias for it. an example:

```
using Index = std::ptrdiff_t

for (Index index { 0 }; index < static_cast<Index>(arr.size()); ++index)
```

4. in cases where you can deduce the type of your loop variable from from the initializer, you can use the `auto` to have the compiler deduce the type.

```
for (auto index {static_cast<std::ptrdiff_t>( arr.size-1 )}; index >= 0; --index)
```

in the c++23 the `z` suffix was introduced as a literal definer for the signed counterpart of `std::size_t` (probably `std::ptrdiff_t`)

```
for (auto index {0Z}; index < static_cast<std::ptrdiff_t>(arr.size()); ++index)
```

Getting the length of an array as a signed value

1. using `static_cast` to cast the return of `std::size()` or `.size()` to a signed. an example:

```
#include <vector>
#include <iostream>

using Index = std::ptrdiff_t

int main()
{
    std::vector arr { 1, 2, 3, 4, 5 };

    for (auto index { 0Z }; index < static_cast<Index>(arr.size()-1); ++index)
        std::cout << arr[static_cast<std::size_t>(index)] << ' ';
}
```

this works well but the downside is how much it clutters code. one way to eliminate this is removing the length variable out of the loop block

```
#include <iostream>
#include <vector>

using Index = std::ptrdiff_t

int main()
{
    std::vector arr { 1, 2, 3, 4, 5 };

    auto len { static_cast<Index>(arr.size) };
    for (auto index { len-1 }; index >= 0; --index)
        std::cout << arr[static_cast<std::size_t>(index)] << ' ';
}
```

more evidence that the designers of c++ want to push the paradigm of interacting with arrays and their indices using signed integers is the introduction of `std::ssize()` in C++20. the function returns the size of a container as a signed integer (likely `std::ptrdiff_t`)

```
#include <iostream>
#include <vector>

int main()
{
    std::vector vec { 1, 2, 3, 4, 5 };

    for (auto index{ std::ssize(vec)-1 }; index >= 0; --index)
        std::cout << vec[static_cast<std::size_t>(index)] << ' ';
}
```

Converting the signed loop variable to an unsigned index

once we have a signed loop variable, we would run into warning for implicit signed to unsigned conversions. so we need ways to curb this.

solutions:

1. using `static_cast` to cast to unsigned where necessary: works but clutters code.
2. use a conversion function with a short name.

```

#include <iostream>
#include <vector>

using Index = std::ptrdiff_t;

constexpr std::size_t getU( Index index )
{
    return static_cast<std::size_t>(index)
}

int main()
{
    std::vector vec {1, 2, 4, 5, 6, 7};

    auto len { static_cast<Index>(vec.size()) };
    for(Index index { len-1 }; index >= 0; --index)
        std::cout << vec[getU(index)];
}

```

3. creates a custome view for arrays that do exactly what you want.

our custom array wsuld be able to do this 2 to any container that supports indexing

1. retrieve the element of an array.
2. return the size of an array.

array_view.h:

```

#ifndef CUSTOM_ARRAY_VIEW
#define CUSTOM_ARRAY_VIEW

template <typename T>
class s_array_view
{
private:
    T& m_array;

public:
    using Index = std::ptrdiff_t;

    s_array_view(T& array) : m_array { array } {}

    constexpr auto& operator[](Index index) {
        return m_array[static_cast<typename T::size_type>(index)];
    }
}

```

```
constexpr const auto& operator[](Index index) const {
    return m_array[static_cast<typename T::size_type>(index)];
}
constexpr auto ssize() {
    return static_cast<Index>(m_array.size());
}
};
#endif
```

main.cpp:

```
#include <iostream>
#include <vector>
#include "array_view.h"

int main()
{
    std::vector array { 1, 2, 3, 4, 5, 6 };
    s_array_view sv_array { array };

    for (auto index { sv_array.ssize()-1 }; index >= 0; --index)
        std::cout << sv_array[index] << ' ';
}
```

The only sane choice: avoid indexing altogether!

all options presented here has their own downsides. so its difficult to recommend a best practice.

however theres a choice far more sane than the others: avoid indexing with integral values altogether.

this can be done using for-each loops or iterators. they will be discussed later.

and we should use them than interacting with intergral values for traversals whenever possible (*best practice*).

Range-based for loops (for-each)

altho forloops provide in interface for traversing through a container, they often lead to one-off errors, easy to mess up and are subjected to array indexing sign problems.

because traversing forward in an array is such a common thing to do c++ provides for loop called the `ranged based` for loop. which helps iteration and traversal of a container without explicit indexing.

range based for loops are simpler, safer and work with all common array types in c++ including `std::vector`, `std::array` and c-style arrays.

Range-based for loops

syntax:

```
for (element_declaration : array_object)
    statement;
```

when a range based loop is encountered, the loop is traversed and assigns the current iterated upon element to the `element_declaration`, then the statement is executed.

for best result `element_declaration` should have the same type as the elements in the array or else type conversion will occur.

an example:

```
#include <iostream>
#include <vector>

int main()
{
    std::vector fibonacci { 0, 1, 1, 2, 3, 5, 8, 13, 21, 34, 55, 89 };

    for (int num : fibonacci)
        std::cout << num << ' ';
    std::cout << '\n';
}
```

because `num` is assigned the value of the elements of the array, through copy assignment this can be quite expensive for some datatypes.

Range-based for loops and type deduction using the `auto` keyword

because `element_declaration` should be of the same data-type as the elements in the array to be traversed. this is an ideal case to use the `auto` keyword for type deduction of the element.

an example:

```
#include <vector>
#include <iostream>
```

```
int main()
{
    std::vector array { 1, 2, 3, 4, 5, 6 };

    for (auto num : array)
        std::cout << num;
}
```

this is a best practice.

Avoid element copies by using references

copying some elements in an array can be expensive so we can assign them to references thereby circumventing the problem of the making expensive copies.

```
#include <vector>
#include <iostream>
#include <string>

int main()
{
    using namespace std::literals;
    std::vector words{ "peter"s, "likes"s, "frozen"s, "yogurt"s };

    for (const auto& word : words)
        std::cout << word << ' ';
    std::cout << '\n';
}
```

if the element_declaration is made non-const we will be able to modify the elements of the array(something not possible if our element declaration is a copy of the value).

Consider always using const auto& when you don't want to work with copies of elements

normaly we use:

- `auto` : for cheap to copy elements
- `const auto&` : for elements that are expensive to copy and we require only read access to them.
- `auto&` : for elements that are expensive to copy but we desire both read and write access to them.

case where using auto can be acceptable:

```
using namespace std::literals;
std::vector words{ "peter"sv, "likes"sv, "frozen"sv, "yogurt"sv };
// sv is a literal suffix string_views.

for (auto word : words) // the auto keyword is fine here only because the
                          // type deduces to a string_view that acts as an inexpensive view for strings
                          // (just like a reference but only for strings).
    std::cout << word;
```

for expensive to copy elements like strings: a string_view, const auto& or auto& are more memory conservative.

range based for loops don't work on decayed C-style arrays (whatever the fuck that is).

Getting the index of the current element

range based for loops don't provide an interface to access the index of the elements it traverses over. but this issue can be resolved by creating your own personal counter.

Range-based for loops in reverse : c++20

range based for loops only traverse in a forward order, however there are cases where we might want to traverse an array in a reverse order.

in c++20 we can use the std::views::reverse of the <ranges> library to create a reverse view of the elements to be traversed.

an example:

```
#include <iostream>
#include <vector>
#include <ranges> // c++20
#include <string_view>

int main()
{
    using namespace std::literals;
    std::vector words{ "Alex"sv, "Bobby"sv, "Chad"sv, "Dave"sv };

    for (const auto& word : std::views::reverse(words))
        std::cout << word;
}
```


Array indexing and length using enumerators

one of the biggest flaws of array indexing is that the subscript tells us nothing about the index or the element to be retrieved.

using unscoped enumerators for indexing

subscripts of containers like `std::vector` and its likes have their type as `container::size_type` which is usually an alias for `std::size_t`.

and since unscoped enumerations will implicitly convert to `std::size_t` we can use the for indexing to help document the meaning of an index.

```
#include <vector>

namespace Students
{
    enum Names
    {
        toby,
        dwight,
        jan,
        michael,
        creed,
        kevin,
        max_employees
    };
}

int main()
{
    std::vector scores { 22, 333, 455, 666, 88, 999 };
    scores[Students::toby] = 100;
    // here were updating the score of student toby.
}
```

because enumerations are implicitly constexpr, conversion of an enumerator to an unsigned integral type is not considered a narrow conversion, thus avoiding signed/unsigned indexing problems.

Using a non-constexpr unscoped enumeration for indexing

the underlying type of enumerators are implicitly defined (so they could be either signed or unsigned). and because they are implicitly constexpr we won't run into any warnings when using them to index arrays directly (this is because constexprs do not narrow convert if they are used in a context that can work but might narrow convert for normal variables types: I suspect this is due to compiler optimizations for compile time variables like constexprs).

so we experience no narrow conversion warnings when used directly but we might experience if the enumerator is made non-constexpr:
an example:

```
#include <vector>

namespace DunderMifflin
{
    enum Scranton
    {
        toby,
        dwight,
        jan,
        michael,
        creed,
        kevin,
        max_employees
    };
}

int main()
{
    std::vector workScores {7, 9, 9, 5, 6, 7};
    DunderMifflin::Scranton boss { DunderMifflin::michael };           // non-constexpr

    workScores[boss] = 10;                                             //
    // possible narrow conversion warning because boss is now a non constexpr
    // enumerator,
    // and
    // might default to a signed underlying type.
}
```

we can fix this by specifying the underlying type of the enumerations.
an example:

```

#include <vector>

namespace DunderMifflin
{
    enum Scranton : unsigned int    // specified.
    {
        toby,
        dwight,
        jan,
        micheal,
        creed,
        kevin,
        max_employees
    };
}

int main()
{
    std::vector workScores {7, 9 ,9, 5, 6, 7};
    DunderMifflin::Scranton boss { DunderMifflin::micheal };           // non-
    constexpr but unsigned,

    workScores[boss] = 10;                                             // sign
    conversion does not occur since unsigned int can be converted to std::size_t
    with no issues

                                                                    // i
    think that integral promotion?
}

```

Using a count enumerator

if you notice we've been adding the enumeration `max_employees` as the last enumerator to all our enumerators. if all prior enumerators of the enumeration are using their default values this means that the `max_employees` enumerator should hold length value or count value of the valid enumerators.

we call this the `count enumerator` as it represents the count of the prior enumerators: an example:

```

#include <vector>

namespace DunderMifflin
{
    enum Scranton : unsigned int    // specified.
    {

```

```

        toby,
        dwight,
        jan,
        micheal,
        creed,
        kevin,
        max_employees
    };
}

int main()
{
    std::vector<int> scores (DunderMifflin::max_employees);           //
    create a vector that holds 5 elements.
    scores[DunderMifflin::kevin] = 78;                               //
    updating the score that belongs to kevin.
    std::cout << DunderMifflin::max_employees << '\n';             // here
    its used to prints the length of the scores or how many employees there are.
}

```

this technique is nice because if an enumerator should get larger we have the nice interface for adding new enumerators and the length is automatically updated as long as any new enumerator is declared before the count enumerator making it always last to hold the count.

Asserting on array length with a count enumerator

this is useful for always making sure that the enumeration count enumerator is of the right value: which should be equal to the size of the vector it acts as an indexer for.
an example:

```

#include <cassert>
#include <iostream>
#include <vector>

enum StudentNames
{
    kenny, // 0
    kyle, // 1
    stan, // 2
    butters, // 3
    cartman, // 4
    max_students // 5
};

```

```

int main()
{
    std::vector scores { 1, 2, 7, 4, 5, 6, 7, 8 };

    assert(std::size(scores) == max_students);
    // proceed to interact with array using its mapped unscoped enumerations
    interface for indexing.
}

```

use `static_assert` if your array is constexpr, vectors do not support constexprs but `std::array` does.

Arrays and enum classes

because unscoped enumerators pollute the scope they are declared, we can opt to use the scoped enumerators: `enum classes`

but they do not support implicit conversions to their underlying types and because of this might run into issues when trying to use their enumerators as array indices.

```

#include <vector>
#include <iostream>

enum class StudentNames // now an enum class
{
    kenny, // 0
    kyle, // 1
    stan, // 2
    butters, // 3
    cartman, // 4
    max_students // 5
};

int main()
{
    std::vector<int> students ( static_cast<int>
(StudentNames::max_students) );

    students[static_cast<size_t>(StudentNames::stan)] = 33;
    std::cout << static_cast<int>(StudentNames::max_student) << '\n';
}

```

although this works in eradicating the pollution of the current scope problem, it also introduces some new ones:

1. a pain to write, clutters code a lot due to the need to specify the namespace.
2. a lot of `static_cast` which leads to more clutters
3. due to all the wrong might lead to errors.

we can fix this though by overloading the `operator+`:

```
#include <vector>
#include <iostream>

enum class StudentNames // now an enum class
{
    kenny, // 0
    kyle, // 1
    stan, // 2
    butters, // 3
    cartman, // 4
    max_students // 5
};

constexpr auto operator+(StudentNames a) noexcept {
    return static_cast<std::underlying_type_t<StudentNames>>(a);
}

int main()
{
    std::vector<int> students (+StudentNames::max_students);

    students[static_cast<size_t>(StudentNames::stan)] = 33;
    std::cout << +StudentNames::max_students << '\n';
}
```

std::vector resizing and capacity

in this subtopic we are focusing on what makes `std::vector` significantly different than other array types: the ability to resize itself after being instantiated.

Fixed-size arrays vs dynamic arrays

most arrays have a limitation that is the size should be known at the point of instantiation and then cannot be changed. such arrays are called `fixed-size arrays`.

both `std::array` and `C Style arrays` are fixed lengthed.

on the other hand `std::vector` is a `dynamic array`. a dynamic array is an array whose size can be changed after instantiation. the ability to be resizable is what makes `std::vector` special.

Resizing a `std::vector` at runtime

this is done by calling the `.resize()` with the size as the argument.
an example:

```
#include <vector>
#include <iostream>

int main()
{
    std::vector vec { 1, 2, 4 };
    std::cout << "the length is: " << vec.size() << '\n';

    vec.resize(5);
    std::cout << "the length is: " << vec.size() << '\n';

    for (auto i : vec)
        std::cout << i << ' ';
    std::cout << '\n';
}
```

this prints

```
the length is: 3
the length is: 5
1 2 4 0 0
```

vectors can also be resized to be smaller: it preserves the first occurring elements and discards the later elements to fit the new size.

The length vs capacity of a `std::vector`

the length of a vector is the number of elements in a vector that is considered in use.

the *capacity* is the number of elements with allocated memory in the vector: it comprises of both in use and not in use elements.

we can get the capacity of a vector by calling the `.capacity()` function.

Reallocation of storage, and why it is expensive

when `std::vector` changes the amount of storage its managing this is known as **reallocation**.

what happens under the hood when reallocation occurs:

1. `std::vector` acquires new memory sufficient for the new capacity of element to be catered for
2. the elements in the old memory gets copied (or moved, if possible) to the newly allocated storage and the old memory gets returned back to the system.
3. the capacity and length of the `std::vector` object are updated to conform to the new storage.

what happens when reallocation occurs is basically hidden replacement with the old `std::vector` object being replaced by a new one.

sounds expensive? thats because it is. we should avoid resizing where possible.

Why differentiate length and capacity?

basically at a vectors initialization capacity and length are the same.

when the vector gets expanded (via resizing and reallocation) the capacity and length, although incremented, resolve to be equal.

but something special happens when a vector is shrunk.

refer to cpp file : [test_capacity_length_vector.cpp](#)

and run executable file : [test_capacity_length_vector](#)

to notice the difference of them

something to note (personal findings):

1. once the capacity of a vector increases it doesn't reduce (implicitly i think)
2. the length can increase or reduce

tracking capacity separately from length allows `std::vector` to avoid reallocations when length is changed.

Vector indexing is based on length, not capacity

that makes sense: because length keeps track of elements in use.

why would we want to index for an element that isn't in use?

A subscript is only valid if it is between 0 and the vector's length (not its capacity)!

Shrinking a `std::vector`

resizing a vector to be larger will increase the vectors length, and will increase the capacity if required. however resizing a vector to be smaller will only decrease its length, and not its capacity.

to fix this vectors have a member function `.shrink_to_fit()` that request that the vector shrinks its capacity to match its length

the requests is non-binding meaning the compiler is free to ignore it.

an example:

```
#include <vector>
#include <iostream>

void printCapLen(const std::vector<int>& v)
{
    std::cout << "Capacity: " << v.capacity() << " Length:" << v.size()
<< '\n';
}

int main()
{
    std::vector num { 1, 3, 4, 5, 6 };
    printCapLen(num);

    num.resize(0);
    printCapLen(num);

    num.shrink_to_fit();
    printCapLen(num);
}
```

when `shrink_to_fit()` is called the compiler reallocated its capacity to 0: freeing 1000 allocated objects.

std::vector and stack behavior